




AIMS

African Institute for
Mathematical Sciences
RWANDA



AMMI Review sessions

Deep Learning (2) Machine Learning Basics (2)



Performance evaluation metrics

- Choosing the right metric is crucial while evaluating machine learning (ML) models
- Various metrics are proposed to evaluate ML models in different applications

Performance evaluation metrics

- Choosing the right metric is crucial while evaluating machine learning (ML) models
- Various metrics are proposed to evaluate ML models in different applications
- What is the difference between the loss function and a metric ?
 - Loss functions are functions that show a measure of the model performance and are used to train a machine learning model and are usually differentiable in model's parameters.
 - metrics are used to monitor and measure the performance of a model (during training, and test), and do not need to be differentiable
 - In some tasks the performance metric is differentiable, it can be used both as a loss function and a metric, such as MSE.

Popular Performance evaluation metrics

- Popular evaluation metrics includes but not limited to:
 - *Classification Metrics (accuracy, precision, recall, F1-score, ROC, AUC, ...)*
 - *Regression Metrics (MSE, MAE)*
 - *Ranking Metrics (MRR, DCG, NDCG)*
 - *Statistical Metrics (Correlation)*
 - *Computer Vision Metrics (PSNR, SSIM, IoU)*
 - *NLP Metrics (Perplexity, BLEU score)*
 - *Deep Learning Related Metrics (Inception score, Frechet Inception distance)*

Classification Metrics

- Let's first review the **Confusion Matrix** concept:
- Tabular visualization of the model predictions versus the ground-truth labels

- True positive(TP) — Correct positive prediction
- False positive(FP) — Incorrect positive prediction
- True negative(TN) — Correct negative prediction
- False negative(FN) — Incorrect negative prediction

		Actual Values	
		Positive (1)	Negative (0)
Predicted Values	Positive (1)	TP	FP
	Negative (0)	FN	TN

Classification Metrics

- Classification metrics derived from the confusion matrix:

- Accuracy** = $\frac{TP+TN}{P+N}$ how many samples predicted correctly

- Precision**(Positive predicted value) = $\frac{\text{true positives}}{\text{predicted positives}} = \frac{TP}{TP+FP}$

can we trust your positive predictions ?

- Recall/Sensitivity**(True positive rate) = $\frac{\text{true positives}}{\text{actual positives}} = \frac{TP}{TP+FN}$

Of all of the true samples how many of them we were able to predict correctly/ how much we missed ?

		Predicted class	
		P	N
Actual Class	P	True Positives (TP)	False Negatives (FN)
	N	False Positives (FP)	True Negatives (TN)

Sensitivity Recall P (points to TP)
 Specificity N (points to TN)
 Precision (points to TP and FP)

Classification Metrics

- Specificity (True negative rate) = $\frac{TN}{N}$
- F1 Score(Harmonic mean of precision and recall) =
$$\frac{(1+b)*Precision*Recall}{b^2Precision+Recall}$$
 where b is commonly 0.5, 1, 2
- Depending on application, you may want to give higher priority to recall or precision
- In medical diagnosis for example you try to avoid false negatives → requires high recall
- For a spam binary classifier you try to avoid false positives → requires high precision

		Predicted class	
		P	N
Actual Class	P	True Positives (TP)	False Negatives (FN)
	N	False Positives (FP)	True Negatives (TN)

Annotations:

- Sensitivity / Recall / P:** Indicated by a yellow box around the TP and FN cells.
- Specificity / N:** Indicated by a green box around the FP and TN cells.
- Precision:** Indicated by a red box around the TP and FP cells.

Classification Metrics

- The confusion Matrix for Multi-class
- Let's say you have N classes, then your confusion matrix would be an $N \times N$ matrix
- With the left axis showing the true class, and the top axis showing the class assigned to an item with that true class.
- Each element i, j of the matrix would be the number of items with true class i that were classified as being in class j .
- You then can calculate the precision recall per class, how ?! (exercise)

[illegible]

ROC Curve



- The **receiver operating characteristic curve** is plot which shows the performance of a binary classifier as function of its cut-off threshold.
- It essentially shows the true positive rate (TPR) against the false positive rate (FPR) for various threshold values.
- As an example your model may predict the below probabilities for 4 sample images: [0.45, 0.6, 0.7, 0.3].
- Then depending on the threshold values below, you will get different labels:
 - cut-off= 0.5: predicted-labels= [0,1,1,0] (default threshold)
 - cut-off= 0.2: predicted-labels= [1,1,1,1]
 - cut-off= 0.8: predicted-labels= [0,0,0,0]

ROC Curve

- For each of these 4 different thresholds we can construct a confusion matrix and calculate the TRP and FPR.
- ROC curve essentially finds out the TPR and FPR for various threshold values and plots TPR against the FPR

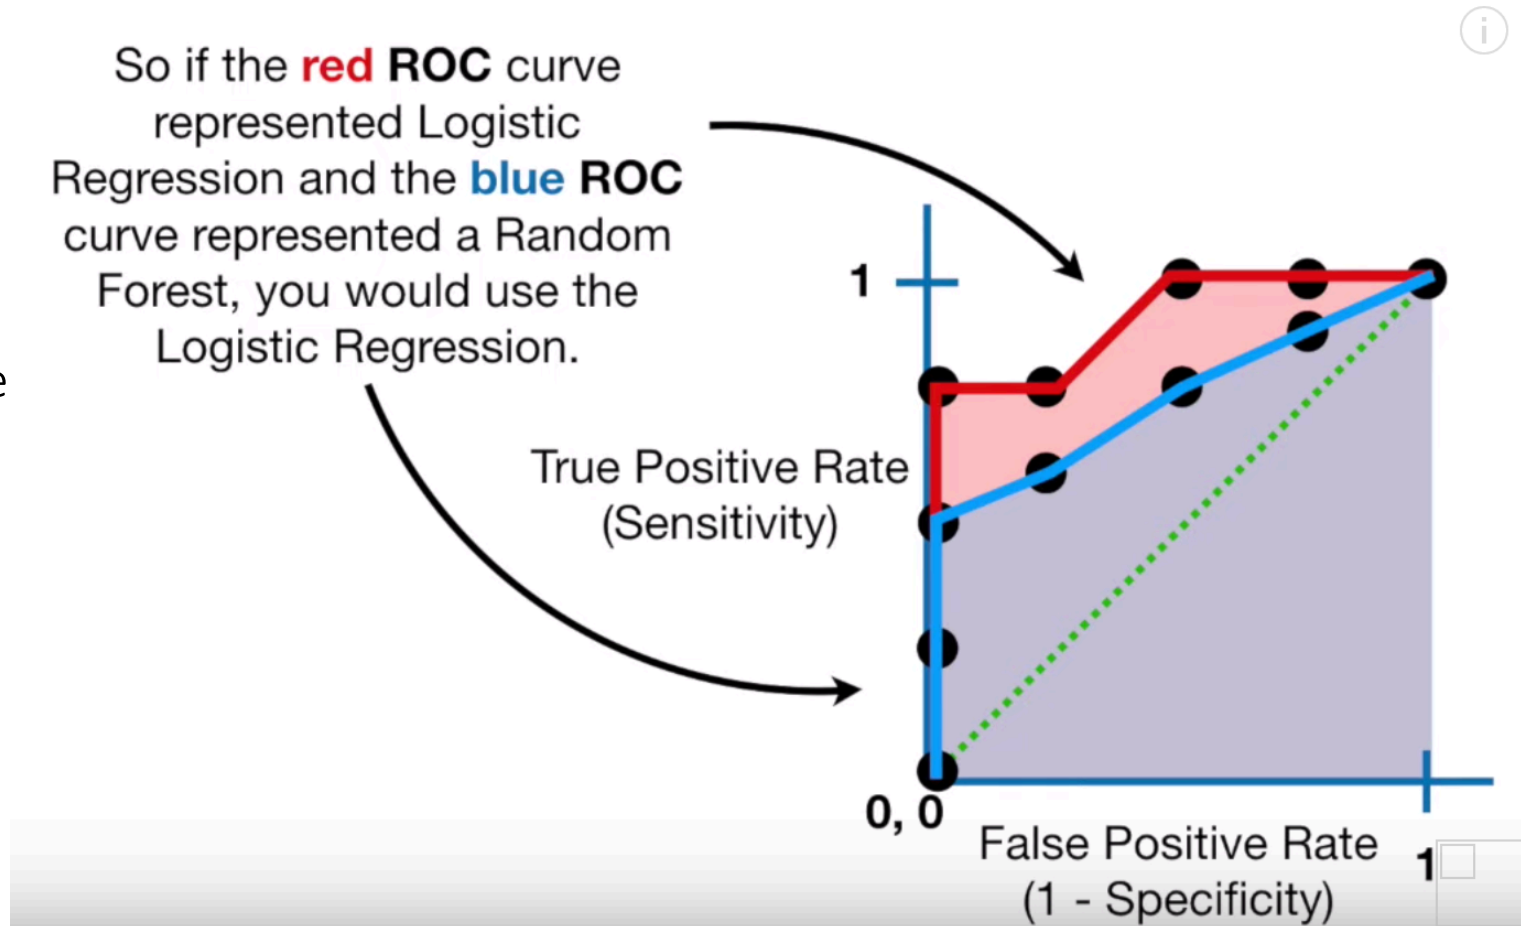
ROC Curve

- The lower the cut-off threshold on positive class, the more samples predicted as positive class, i.e. higher true positive rate (recall) and also higher false positive rate
- Therefore, there is a trade-off between how high the recall could be versus how much we want to bound the error (FPR).
- ROC curve is a popular curve to look at overall model performance and pick a good cut-off threshold for the model.
- [Detailed example](#)



Area Under The Curve (AUC)


- An aggregated measure of performance of a binary classifier on all possible threshold values
- Can be used to compare the performance of two different models
- if your application requires higher recall then you might choose a model using a point the ROC curve even if the AUC value is not too high



Regression Related Metrics

Mean squared error	$\text{MSE} = \frac{1}{n} \sum_{t=1}^n e_t^2$
Root mean squared error	$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{t=1}^n e_t^2}$
Mean absolute error	$\text{MAE} = \frac{1}{n} \sum_{t=1}^n e_t $

The MAE is more robust than the MSE
(less sensitive to outliers and fluctuations)




Pearson's Correlation Coefficient

- Correlation is a technique for investigating the relationship between two quantitative, continuous variables, for example, age and blood pressure.
- Pearson's correlation coefficient (r) is a **measure of the strength of the association** between the two variables.
- The first step in studying the relationship between two continuous variables is to draw a scatter plot of the variables to check for linearity.
- The nearer the scatter of points is to a straight line, the higher the strength of association between the variables.

Pearson's Correlation Coefficient


- $r = -1$
- data lie on a perfect straight line with a negative slope



- $r = 0$
- no linear relationship between the variables



- $r = +1$
- data lie on a perfect straight line with a positive slope



Cross validation


Split		Exp. A	Exp. B	Exp. C
Train	Valid	Performance A1	Performance B1	Performance C1



Test

K-fold cross-validation

Splits	Exp. A	Exp. B	Exp. C
<div><div>Train</div><div>Valid</div></div>	Performance A_1	Performance B_1	Performance C_1
<div><div></div><div></div><div></div></div>	Performance A_2	Performance B_2	Performance C_2
<div><div></div><div></div></div>	Performance A_K	Performance B_K	Performance C_K
	Average perf. A StdError perf. A	Average perf. B StdError perf. B	Average perf. C StdError perf. C



K-fold cross-validation

Splits	Exp. A	Exp. B	Exp. C
<div><div>Train</div><div>Valid</div></div>	Performance A_1	Performance B_1	Performance C_1
<div><div></div><div></div><div></div></div>	Performance A_2	Performance B_2	Performance C_2
<div><div></div><div></div></div>	Performance A_K	Performance B_K	Performance C_K
	Average perf. A StdError perf. A	Average perf. B StdError perf. B	Average perf. C StdError perf. C


Test

- K-fold is less biased but it's computationally expensive, in the above example we train the model 9 times and validate it 9 times and test 1 time!

Normalization


- Feature scaling

$$x' = \frac{x - x_{min}}{x_{max} - x_{min}}$$



- Standardization



$$z = \frac{x - \mu}{\sigma}$$



- In scaling, you're changing the range of your data while in normalization you're mostly changing the shape of the distribution of your data.
- You need to normalize our data if you're going to use a machine learning or statistics technique that assumes that data is normally distributed

Hyperparameter search



Random search with good candidates is already a strong baseline.



Source: Bergstra, James, and Yoshua Bengio. "Random search for hyper-parameter optimization." Journal of Machine Learning Research 13, no. Feb (2012): 281-305.

Hyperparameter search



Random search with good candidates is already a strong baseline.



Source: Bergstra, James, and Yoshua Bengio. "Random search for hyper-parameter optimization." Journal of Machine Learning Research 13, no. Feb (2012): 281-305.

Hyperparameter search

Random search with good candidates is already a strong baseline.



Source: Bergstra, James, and Yoshua Bengio. "Random search for hyper-parameter optimization." Journal of Machine Learning Research 13, no. Feb (2012): 281-305.

References

- Ian Goodfellow, Yoshua Bengio, and Aaron Courville. “**Deep learning**”, MIT press, 2016.
- [Shervin Minaee 20 Popular Machine Learning Metrics](#)
- [Ivado-mila Deep learning summer school 2019](#)