




AIMS

African Institute for
Mathematical Sciences
RWANDA




AMMI Review sessions

Deep Learning (5) Convnets II



Parameter sharing

- Parameter sharing refers to using the same parameter for more than one function in a model, i.e the value of the weight applied to one input is tied to the value of a weight applied elsewhere
- kernels can extract similar features at multiple places in an image
- Convolutions are **equivariant to translation** - i.e. translating an input in space will result in a translated convolution output
- Convolution is not naturally equivariant to some other transformations, such as changes in the scale or rotation of an image.
- Useful features which have been learned can be **reused** at multiple regions in an image



Convolution arithmetics

To Calculate the output dimensions of an input ($W \times H$) and a filter with the size (F_w, F_h), horizontal stride S_w and a vertical one S_h with Padding P :

- Output width = $\frac{(W - F_w + 2P)}{S_w} + 1$
- Output height = $\frac{(H - F_h + 2P)}{S_h} + 1$


To calculate the output dimensions of pooling operation

- Output width = $\frac{(W - F)}{S} + 1$
- Output height = $\frac{(H - F)}{S} + 1$

Batch Normalization

Z-score input: $\hat{\mathbf{x}} = \frac{\mathbf{x} - \mu}{\sigma}$

Affine transform: $\mathbf{y} = \gamma \hat{\mathbf{x}} + \beta$




* Note: Batch normalization works differently in training and test time.

Receptive Field

- The **receptive field** of a feature is the part of the input that “influenced” the feature:


Feature maps

Input image



* Credits: Dang Ha The Hien

Receptive Field



Kernel Size (k_1): 3
Padding (p_1, q_1): 1
Stride (s_1): 3

Kernel Size (k_2): 2
Padding (p_2, q_2): 0
Stride (s_2): 1

Receptive Field

Filter sizes

- What is the receptive field of one **5x5 filter**?
 - And how much compute does it take?
- What is the receptive field of **two 3x3 filters**?
 - How much compute does that take?
- Many current architecture use primarily **3x3 filters** for this reason

Receptive Field

- The deeper you go, the wider the receptive field

n : number of features
 r : receptive field size
 j : jump (distance between two consecutive features)
 $start$: center coordinate of the first feature

k : convolution kernel size
 p : convolution padding size
 s : convolution stride size


$$n_{out} = \left\lfloor \frac{n_{in} + 2p - k}{s} \right\rfloor + 1$$

$$j_{out} = j_{in} * s$$


$$r_{out} = r_{in} + (k - 1) * j_{in}$$

$$start_{out} = start_{in} + \left(\frac{k - 1}{2} - p \right) * j_{in}$$


Layer 0: $n_0 = 5$; $r_0 = 1$; $j_0 = 1$;
 $start_0 = 0.5$




Conv1: $k_1 = 3$; $p_1 = 1$; $s_1 = 2$



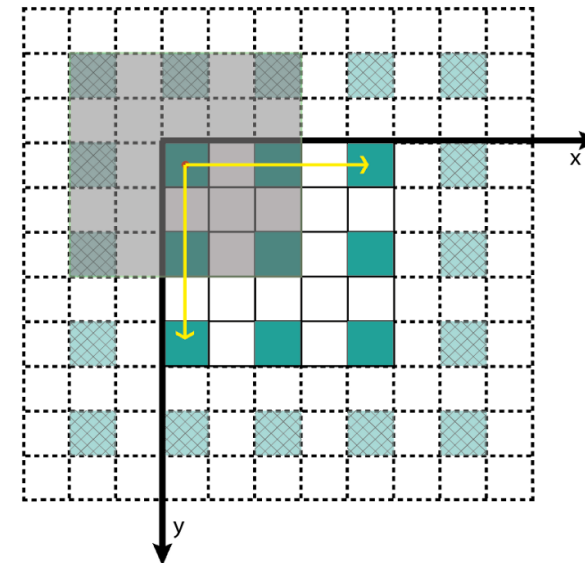
Layer 1: $n_1 = 3$; $r_1 = 3$; $j_1 = 2$;
 $start_1 = 0.5$




Layer 1: $n_1 = 3$; $r_1 = 3$; $j_1 = 2$;
 $start_1 = 0.5$



Conv2: $k_2 = 3$; $p_2 = 1$; $s_2 = 2$




Layer 2: $n_2 = 2$; $r_2 = 7$; $j_2 = 4$;
 $start_2 = 0.5$



Dilated convolution

- Another way to downsample our input is by using **dilated convolutions** (or *atrous* convolution). The **dilation coefficient** d determines pixel spacing between convolutions.



Transposed convolution

- **Transposed convolutions** are used to **increase** the resolution of a **feature map**. It can be thought of as an interpolation, or a convolution on the output. It is often used in the context of **upsampling** the output of a network.
- Convolution as a matrix operation

For example This linear operation takes the input matrix flattened as a 16-dimensional vector and produces a 4-dimensional vector that is later reshaped as the 2×2 output matrix.

$$\begin{pmatrix} w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 & 0 & 0 & 0 & 0 \\ 0 & w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 \\ 0 & 0 & 0 & 0 & 0 & w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} \end{pmatrix}$$

- To get the other way around from 4 dimensions to 16 we can use the transposed matrix
- The gradients w.r.t. the inputs can be computed by “Transposed convolution

References

- [Dang Ha The Hien, A guide to receptive field arithmetic for Convolutional Neural Networks.](#)
- [Vincent Dumoulin, Francesco Visin, A guide to convolution arithmetic for deep learning](#)
- [Ivado-mila Deep learning summer school 2019](#)